

# Sultan Mahmud Sajal

✉ sxs2561@psu.edu

🌐 smsajal.github.io

☎ +1-814-380-3595

## Research Interests

My research focuses on **workload analysis** and **performance evaluation** of **Cloud Computing Systems**. I am especially interested in faithfully scaling, analysing and synthesizing cloud workloads to facilitate realistic experimentation in both academia and industry. Much of my research has explored how different characteristics of the workload affect the performance of the system under different conditions and how we can modify the workload to our need while preserving important characteristics.

## Education

- Expected 2023    📖 **Ph.D. in Computer Science and Engineering, The Pennsylvania State University**  
**Advisors:** *Dr. Timothy Zhu* and *Dr. Bhuvan Urgaonkar*  
**Selected Courses:** Operating Systems Design, Performance Evaluation, Topics in Cloud Computing, Data Mining, Algorithm Design and Analysis, Computer Networks, Programming Language, Natural Language Processing, Database Systems
- 2013 - 2017    📖 **B.Sc. in Computer Science and Engineering, Bangladesh University of Engineering and Technology**  
**Thesis topic:** An Empirical Study on the Growth of New Languages and Their Users in Stack Overflow  
**Advisor:** *Dr. Rifat Shahriyar*  
**Selected Courses:** Artificial Intelligence, Machine Learning, Operating System, Computer Networks, Database.

## Publications

- 1 **Sajal, Sultan Mahmud\*** and Hasan\*, Rubaba, Zhu, T., Urgaonkar, B., & Sen, S. (2021). **TraceSplitter: a new paradigm for downscaling traces** [\*Equal Contribution]. *Proceedings of the Sixteenth European Conference on Computer Systems, (EuroSys '21)*, 606–619.
- 2 **Sajal, Sultan Mahmud**, Mehrab, Z., Zaman, I., Uddin, M., & Rahman, A. (2017). **Poster Abstract: Handwriting Recognition Using Accelerometer**. *International Conference on Networking Systems and Security, (NSysS)*.





## Research Experience

- 2020 - Ongoing    📖 **Upscale Workloads from Cloud Infrastructure and Large Datacenters**  
- Mentors: *Dr. Timothy Zhu* and *Dr. Bhuvan Urgaonkar* (The Pennsylvania State University), *Dr. Siddhartha Sen* (Microsoft Research)  
- Upscale real workload for enabling faithful systems experimentation under varying loads  
- Exploring novel approaches to faithfully increase the load of cloud workloads  
- Performance evaluation using 3-tier web server system consisting of Mediawiki web server, Memcached and MySQL database
- May, 2022 - Aug, 2022    📖 **Capacity Reservation in Cloud System**  
- Mentors: *Dr. Luke Marshall*, *Dr. Beibin Li* and *Dr. Ishai Menache* (Cloud Operations Research (CORE) - MSR)  
- Develop capacity reservation techniques for Azure which maximizes sellable capacity while minimizing SLA violation risk  
- Extend existing simulator to analyze different capacity reservation techniques for comparison





## Research Experience (continued)

- May, 2021 - August, 2021     **Development of Flight Simulator for Spark Jobs**  
- Mentors: *Dr. Abhishek Roy* and *Dr. Joyce Cahoon* (Gray Systems Lab (GSL), Microsoft)  
- Create realistic benchmark for Spark Workloads from query traces  
- Generate synthetic representative datasets for the benchmark
- 2018 - 2020     **Downscale Workloads from Cloud Infrastructure and Large Datacenters**  
- Mentors: *Dr. Timothy Zhu* and *Dr. Bhuvan Ugaonkar* (The Pennsylvania State University), *Dr. Siddhartha Sen* (Microsoft Research)  
- Downscale cloud workload while preserving important characteristics to facilitate realistic systems research and industry prototyping  
- Proposed novel techniques for realistically downscaling workloads that preserve important characteristics such as arrival process and performance  
- For performance evaluation, implemented a 3-tier web server system consisting of Elgg web server, Memcached and MySQL database using Docker containers on Linux  
- For real world case study, implemented autoscaler on Wikimedia application with MySQL database  
- Used statistical methods (e.g. energy distance) for data analysis
- 2016-2017     **An Empirical Study on The Growth of New Languages and Their users in Stack Overflow**  
- Mentor: *Dr. Rifat Shahriyar* (Bangladesh University of Engineering and Technology)  
- Studied the growth patterns of new language resources and Users in Stack Overflow  
- Analyzed the Stack Overflow data using Java and MySQL on Linux
-  **Revealing Correlations Over Gene Expressions in Hypoxia-Induced Human Heart Cells**  
- Mentors: *Dr. A.B.M. Alim Al Islam* (Bangladesh University of Engineering and Technology), *Dr. Rameen Shakur* (University of Cambridge)  
- Explored the dependency of genetics on Hypoxia  
- Developed analysis framework using MATLAB and Weka

## Professional Experience



- May, 2022 - August, 2022     **Research Intern** at Cloud Operations Research (CORE), Microsoft Research.
- May, 2021 - August, 2021     **Research Intern** at Gray Systems Lab (GSL), Microsoft.
- August, 2018 - Present     **Graduate Research Assistant** and **Graduate Teaching Assistant** at The Pennsylvania State University.
- October, 2017 - July, 2018     **Junior Software Engineer** at Reve Systems.

## Awards

- 2021     Registration Grant, Sigmetrics 2021
- 2020     Student Grant, USENIX Symposium on Operating Systems Design and Implementation (OSDI)
- 2019, 2018     Student Travel Grant, ACM Symposium on Cloud Computing (SoCC)
- 2018     Technical Scholarship, Bangladesh University of Engineering and Technology






## Awards (continued)

---

- 2012     Higher Secondary Certificate Scholarship, Dhaka Education Board, Bangladesh
- 2009     Junior Scholarship, Dhaka Education Board, Bangladesh

## Skills


---

Programming Languages	 Java, Python, C++, Scala, C, R, Scheme, Assembly (Intel 8086)
Analysis Tools	 MATLAB, Weka
Databases	 Oracle, MySQL, SQLite
Frameworks	 Apache Spark, Android, Spring-Boot, JPA, JSP, Scikit-Learn, Bootstrap
Technologies	 AWS Services and SDK, Azure Services and CLI, Docker, Git


## References

---


### Dr. Timothy Zhu

Assistant Professor, Department of Computer Science and Engineering  
The Pennsylvania State University  
 [timothy@cse.psu.edu](mailto:timothy@cse.psu.edu)

### Dr. Bhuvan Urgaonkar

Professor, Department of Computer Science and Engineering  
The Pennsylvania State University  
 [buu1@psu.edu](mailto:buu1@psu.edu)

### Dr. Siddhartha Sen

Principal Researcher  
Microsoft Research in New York City  
 [sidsen@microsoft.com](mailto:sidsen@microsoft.com)